Learning under distribution shift

A synthetic research reading for RUWithMe

1. The question

How does an estimator behave when the test distribution changes?

This working note is invented for product demonstration.

It does not describe a published paper or a validated result.

2. A small experiment

Compare empirical risk minimization with a weighted baseline.

Use three shift strengths and five random seeds.

Record the mean loss and the variability across seeds.

Reading notes and next steps

A synthetic research reading for RUWithMe

3. Questions to carry forward

Which assumption justifies the weight bound?

Does increasing the ambiguity radius improve worst-group loss?

Check sensitivity before drawing any conclusion.

Space for your observations - no progress is inferred from opening this file.